ChatGPT

Updated model-selection map for orchestrated coding tasks (June 18 2026)

Methodology

In response to the user's request to update the previous model-selection map, I re-examined the availability and pricing of models across all four target environments—**OpenAI Codex**, **Google Antigravity IDE**, **Cursor**, and **VS Code/GitHub Copilot**—as of **18 June 2026**. The previous report included retired GPT-5.1 and GPT-5.2 models; those models are now either deprecated or no longer available [1]. I used up-to-date official documentation, pricing tables and third-party analyses to compile the current model lineup. Prices are quoted **per 1 million tokens** (input/output combined) because both platform billing and credit conversion use these units. The **average cost** below is the mean of the input and output price for each model (cached-input costs are lower and not included for simplicity). A **cost-rating** (½ star to 5 stars) ranks models from cheapest to most expensive within each platform.

OpenAI Codex (current models)

OpenAI's Codex surface now uses the latest GPT-5 series models. According to OpenAI's pricing page and Verdent's guidance [2] [3], the relevant models are:

- **GPT-5.4-nano** – smallest and cheapest variant. Optimised for speed-critical agent tasks with a 400 k token context window; ideal for quick code completions and small patches. Costs **$0.20 in / $1.25 out** per 1 M tokens [2].
- **GPT-5.4-mini** – budget version of GPT-5.4 with a 200 k context window. Good for sub-agents, interactive edits and everyday coding tasks. Costs **$0.75 in / $4.50 out** [2].
- **GPT-5.4** – full-size model (1 M context). Incorporates GPT-5.3-Codex's coding capability into a general model [3]. Suitable for large-scale refactoring, multi-file planning and complex agent workflows. Costs **$2.50 in / $15 out** [2].
- **GPT-5.5** – flagship model with a 1 M context window; default for ChatGPT-authenticated Codex sessions [4]. Recommended for the most challenging multi-step tasks, but expensive: **$5 in / $30 out** [2].
- **GPT-5.3-Codex** – specialised coding model still available via API; recommended for API-key workflows because GPT-5.5 has not yet rolled out to API keys [5]. Supports reasoning-effort settings and runs 25 % faster than GPT-5.2-Codex [6]. Costs **$1.75 in / $14 out** [7].

(Retired models: GPT-5.2-Codex and GPT-5.1-Codex are deprecated and have been removed from ChatGPT sign-in; they remain in the API catalogue only for legacy workflows [1].)

Codex model summary

Model	Tasks (keywords)	Avg. cost (USD/M tokens)	Cost rating
GPT-5.4-nano	quick patches, small completions, sub-agent calls	**$0.73**	★½
GPT-5.4-mini	general coding, interactive edits, small refactors	**$2.63**	★★
GPT-5.3-Codex	coding-centred agents, long running tool chains, 200 k context, reasoning settings	**$7.88**	★★★½
GPT-5.4	large refactors, cross-file planning, system design	**$8.75**	★★★★
GPT-5.5	frontier model; deep reasoning, 1 M context, multi-agent orchestration	**$17.50**	★★★★½
GPT-5.4-pro / GPT-5.5-pro	long-context professional tiers with 180 $/M-token output; used only for enterprise customers [2]	**105**	★★★★★

Google Antigravity IDE (Gemini models)

Antigravity 2.0 is now the unified harness for Google's Gemini models. Pricing information from LLM-Stats and Metacto summarises the current lineup [8] [9]. Antigravity defaults to **Gemini 3.5 Flash**, a model released on 19 May 2026 that scores highly on agentic coding benchmarks and costs **$1.50 in / $9 out** per 1 M tokens [10]. The next-generation **Gemini 3.5 Pro** remains in limited preview and is not generally available [11], so it is omitted from this map.

Key models and their purposes:

- **2.5 Flash-Lite** – ultra-cheap tier; good for lightweight completions and file summaries. $0.10 in / $0.40 out per 1 M tokens [12].
- **3.1 Flash-Lite** – cost-efficient upgrade; recommended for small refactors and retrieval-augmented classification. $0.25 in / $1.50 out [13].
- **2.5 Flash** – balanced speed/cost; good for everyday coding with moderate context. $0.30 in / $2.50 out [14].
- **3 Flash** – "Flash" tier preceding 3.5; still available in some regions. $0.50 in / $3 out [15].
- **3.5 Flash** – default Antigravity model. Excels at multi-step tool use, sub-agent orchestration and long-horizon agent tasks [16]. $1.50 in / $9 out [10].
- **2.5 Pro (≤200 k)** – prior flagship; deeper reasoning and compaction; context up to 1 M tokens; $1.25 in / $10 out [17].
- **2.5 Pro (>200 k)** – same model but charged at higher rate for large contexts; $2.50 in / $15 out [17].
- **3.1 Pro (≤200 k)** – current flagship reasoning model; long context (up to 2 M tokens). $2 in / $12 out [18].
- **3.1 Pro (>200 k)** – extended context pricing; $4 in / $18 out [18].

Antigravity model summary

Model	Tasks (keywords)	Avg. cost (USD/M tokens)	Cost rating
2.5 Flash-Lite	tiny completions, simple summaries	**$0.25**	★½
3.1 Flash-Lite	small refactors, retrieval, chat	**$0.88**	★★
2.5 Flash	general coding, debugging, everyday edits	**$1.40**	★★½
3 Flash	balanced throughput vs cost; legacy but still used	**$1.75**	★★½
3.5 Flash	default agentic model; multi-agent orchestration, long-horizon tasks [16]	**$5.25**	★★★★
2.5 Pro (≤200 k)	deeper reasoning, complex analysis	**$5.63**	★★★★
3.1 Pro (≤200 k)	flagship reasoning model, 2 M context; heavy design, audits	**$7.00**	★★★★½
2.5 Pro (>200 k)	long-context reasoning beyond 200 k	**$8.75**	★★★★½
3.1 Pro (>200 k)	highest reasoning/cost tier	**$11.00**	★★★★★
3.5 Pro (preview)	upcoming model; expected to cost around $15 in / $60 out per 1 M tokens according to early reports [19] ; not generally available, so it should be avoided for now.	n/a	—

Cursor IDE

Cursor bundles its own **Composer 2.5** model for everyday tasks and exposes a catalog of external models. The *ofox.ai* guide explains that Composer 2.5 is built on Moonshot's Kimi K2.5 and offers **two tiers**: **Standard** at $0.50 in / $2.50 out per 1 M tokens and **Fast** at $3 in / $15 out [20] . Composer 2.5 Standard is recommended for routine edits; the Fast variant simply trades cost for lower latency [21] . Cursor's model picker also includes a range of other provider models:

- **Grok Code Fast 1 (xAI)** – cost-effective reasoning model with visible chain-of-thought; ideal for agentic coding loops and tool use [22] . $0.20 in / $1.50 out per 1 M tokens.
- **Grok Build 0.1 (xAI)** – fast build-and-compile model; good for generating code artefacts; $1 in / $2 out [23] .
- **Grok 4.3 (xAI)** – 1 M-context Grok model; balanced reasoning; $1.25 in / $2.50 out [24] .
- **Grok 4.20 (xAI)** – more powerful Grok variant; deeper reasoning and planning; $2 in / $6 out [24] .
- **Composer 2.5 Standard** – default Cursor model; mixture-of-experts; good for general coding and summarisation [20] .
- **Composer 2.5 Fast** – same weights on faster hardware; use when latency matters [21] .

- **Anthropic models** – **Claude Haiku 4.5**, **Sonnet 4.6**, **Opus 4.8** and **Fable 5** (May/June 2026 releases). Pricing: Haiku 4.5 $1 in / $5 out; Sonnet 4.6 $3 in / $15 out; Opus 4.8 $5 in / $25 out [25] ; Fable 5 $10 in / $50 out [26] .
- **OpenAI models** – Cursor users can bring their own OpenAI models. Current practical choices include **GPT-5.4-nano**, **GPT-5.4-mini**, **GPT-5.4**, **GPT-5.5** and **GPT-5.3-Codex** (API-key). Pricing matches OpenAI's rate card [2] [7] .
- **Higher-end models** – **Claude Opus 4.8** and **GPT-5.5** deliver the best reasoning but are very expensive. **Claude Fable 5** offers 1 M context and creative reasoning for interactive fiction or narrative generation [26] .

Cursor model summary

Model	Tasks (keywords)	Avg. cost (USD/ M tokens)	Cost rating
Grok Code Fast 1	iterative loops, tool-use, cheap agent runs	**$0.85**	★½
GPT-5.4-nano	tiny tasks, quick completions	**$0.73**	★½
Composer 2.5 Standard	everyday coding, summarisation	**$1.50**	★★
Grok Build 0.1	compiling/build tasks	**$1.50**	★★
Grok 4.3	medium-complex refactors, analysis	**$1.88**	★★½
GPT-5.4-mini	interactive edits, small agents	**$2.63**	★★½
Claude Haiku 4.5	summarisation, explanations	**$3.00**	★★★
Grok 4.20	deeper reasoning, long summaries	**$4.00**	★★★
GPT-5.3-Codex	coding-centric agents, long tool chains	**$7.88**	★★★½
GPT-5.4	large refactors, multi-file planning	**$8.75**	★★★★
Composer 2.5 Fast / Claude Sonnet 4.6	high-speed coding or balanced reasoning	**$9.00**	★★★★
Claude Opus 4.8	high-end reasoning, large migrations	**$15.00**	★★★★½
GPT-5.5	flagship model; complex multi-step agent flows	**$17.50**	★★★★½
Claude Fable 5	mythos-class creative reasoning; 1 M context	**$30.00**	★★★★★

VS Code / GitHub Copilot

GitHub Copilot switched to a token-based **AI Credit** system on June 1 2026. The official pricing table lists the models currently available [27]. Credits convert to US dollars at roughly $0.01 per credit; to simplify comparison, the table below converts credit costs into USD per 1 M tokens (input + output). Copilot models fall into four categories:

- **Low-cost models** – **Grok Code Fast 1**, **GPT-5 mini** and **Raptor mini**. Useful for quick chat, debugging and small tasks. Average costs range from **$0.85 to $1.13** per 1 M tokens [27].
- **Mid-tier models** – **Gemini 3 Flash** and **Claude Haiku 4.5**. Balanced speed and quality for everyday agentic sessions; **$1.75–$3.00** per 1 M tokens [28].
- **Premium models** – **Gemini 2.5 Pro**, **Gemini 3.5 Flash**, **GPT-5.2**, **GPT-5.4**, **Claude Sonnet 4.x** and **MAI-Code-1-Flash**. Provide stronger reasoning and longer context; costs from **$5.25 to $9** per 1 M tokens [29].
- **Frontier models** – **Claude Opus 4.x**, **GPT-5.5** and **Claude Fable 5**. Highest-end reasoning and 1 M context; best for architecture design, security audits and complex agentic workflows. Costs range from **$15 to $30** per 1 M tokens [30].

Copilot model summary

Model	Tasks (keywords)	Avg. cost (USD/ M tokens)	Cost rating
Grok Code Fast 1	cheap coding loops, tool calls	**$0.85**	★½
GPT-5 mini / Raptor mini	general chat, debugging, simple edits	**$1.13**	★★
Gemini 3 Flash	balanced interactive sessions, moderate agents	**$1.75**	★★½
Claude Haiku 4.5	summarisation, explanation, basic refactoring	**$3.00**	★★★
Gemini 2.5 Pro	stronger reasoning, larger context	**$5.63**	★★★½
GPT-5.2	advanced coding, long refactors	**$7.88**	★★★★
Gemini 3.5 Flash	agent-oriented tasks with 1 M context	**$5.25**	★★★½
GPT-5.4	complex planning, multi-file design	**$8.75**	★★★★
Claude Sonnet 4.x	high-quality reasoning, multi-file refactors	**$9.00**	★★★★
Claude Opus 4.x	deep reasoning, system-level design	**$15.00**	★★★★½
GPT-5.5	flagship model; 1 M context; multi-step workflows	**$17.50**	★★★★½
Claude Fable 5	mythos-class creative model; currently unavailable in Copilot [31]	**$30.00**	★★★★★

Note: GitHub's pricing table includes cache-write and cached-input charges; these can reduce costs by up to 90 % in real workloads. Costs above assume no caching for easier comparison.

Recommendations

1. **Use the lowest-cost models for routine tasks.** In any environment, choose the cheapest model when tasks involve quick completions, small bug fixes or file summaries (e.g., GPT-5.4-nano in Codex, Gemini 2.5 Flash-Lite in Antigravity, Grok Code Fast 1 or Composer 2.5 Standard in Cursor, and Grok Code Fast 1 or GPT-5 mini in Copilot).
2. **Select mid-tier models for moderate complexity.** Models like GPT-5.4-mini, Gemini 3 Flash, Composer 2.5 Standard, Grok 4.3 or Haiku 4.5 deliver a good balance of reasoning and price. They handle medium-sized refactors, interactive edits and moderate agentic workflows.
3. **Reserve premium models for complex design and long-context tasks.** Premium options (GPT-5.3-Codex/GPT-5.4, Gemini 3.5 Flash/2.5 Pro, Composer 2.5 Fast, Claude Sonnet) are well-suited for large refactors, multi-file planning, system design reviews and long-running agent flows.
4. **Only use frontier models when necessary.** Frontier models such as GPT-5.5, Claude Opus 4.8 and Claude Fable 5 offer the deepest reasoning and largest context windows but at the highest cost. Deploy them only for architecture design, security audits, cross-modal reasoning or narrative work where cheaper models fail to deliver acceptable quality.
5. **Check model availability and credentials.** Some models (e.g., GPT-5.5 and Claude Fable 5) may only be available through specific surfaces or require ChatGPT/enterprise plans [11] [31]. Verify access before selecting them in an orchestration workflow.

This updated mapping reflects the models and pricing available on **June 18 2026** and replaces the outdated GPT-5.1/5.2 information in the original report.

---

[1] GPT-5.2 and GPT-5.2-Codex deprecated - GitHub Changelog
https://github.blog/changelog/2026-06-05-gpt-5-2-and-gpt-5-2-codex-deprecated/

[2] [7] Pricing | OpenAI API
https://developers.openai.com/api/docs/pricing

[3] [4] [5] [6] GPT-5 Codex Model Names Explained in 2026 - Verdent Guides
https://www.verdent.ai/guides/gpt-5-codex-model-names-explained

[8] [10] [16] Gemini 3.5 Flash: Benchmarks, Pricing, and Complete Specs
https://llm-stats.com/blog/research/gemini-3.5-flash-launch

[9] [12] [13] [14] [15] [17] [18] Gemini API Pricing May 2026: 3.5 Flash, 3.1 Pro, 2.5 Lite
https://www.metacto.com/blogs/the-true-cost-of-google-gemini-a-guide-to-api-pricing-and-integration

[11] Gemini 3.5: frontier intelligence with action
https://blog.google/innovation-and-ai/models-and-research/gemini-models/gemini-3-5/

[19] Google Gemini 3.5 Pro Nears June Launch With 2 Million Token Context And Deep Think Reasoning
https://www.techtimes.com/articles/317919/20260606/google-gemini-35-pro-nears-june-launch-2-million-token-context-deep-think-reasoning.htm

[20] [21] Cursor Composer 2.5: What's New, Best Models, and How to Set It Up
https://ofox.ai/blog/cursor-composer-2-5-setup-guide-2026/

[22] Grok Code Fast 1 - API Pricing & Benchmarks | OpenRouter
https://openrouter.ai/x-ai/grok-code-fast-1

[23] [24] Grok Build 0.1 - API, Specs, Playground & Pricing - Puter Developer
https://developer.puter.com/ai/x-ai/grok-build-0.1/

[25] Claude API Pricing 2026: Opus 4.8, Sonnet 4.6, Haiku 4.5 Costs
https://www.metacto.com/blogs/anthropic-api-pricing-a-full-breakdown-of-costs-and-integration

[26] Introducing Claude Fable 5 and Claude Mythos 5 - Claude API Docs
https://platform.claude.com/docs/en/about-claude/models/introducing-claude-fable-5-and-claude-mythos-5

[27] [28] [29] [30] [31] Models and pricing for GitHub Copilot - GitHub Docs
https://docs.github.com/en/copilot/reference/copilot-billing/models-and-pricing